Linux Time Sensitive Networking

REAL TIME NETWORKING CUT-THROUGH, TIME SYNC, AND PREEMPTION

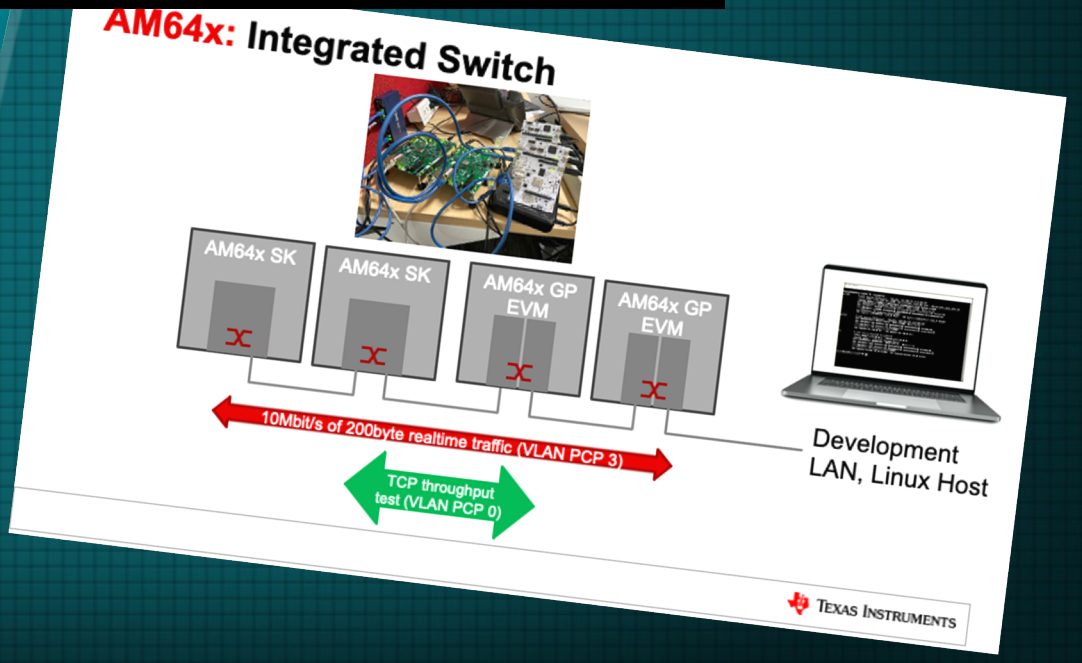

Agenda

- Demo overview
- AM64x TSN features
- Demo
  - Configure cut-through switching with preemption
  - IEEE802.1AS-2011 using ptp4l
  - Express and background traffic
- Q & A

AM64x: Integrated Switch

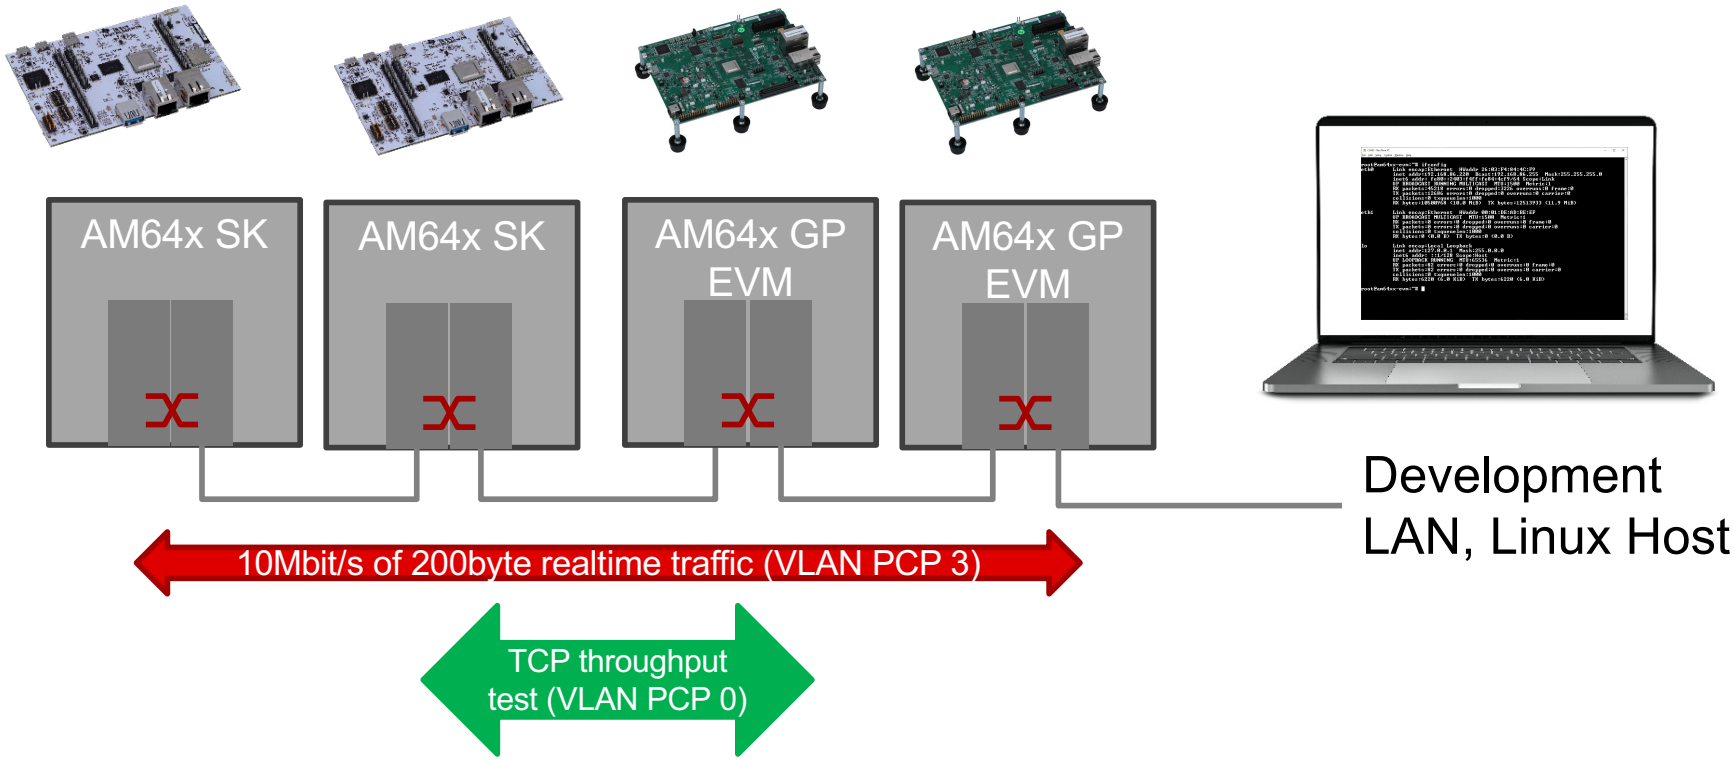

AM64x: Integrated Switch

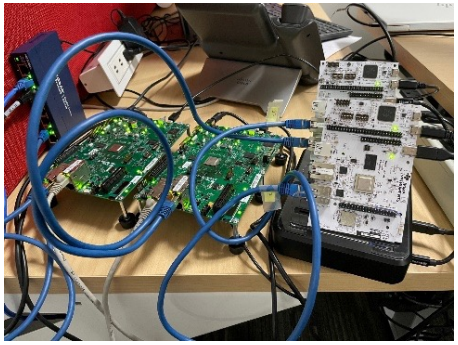

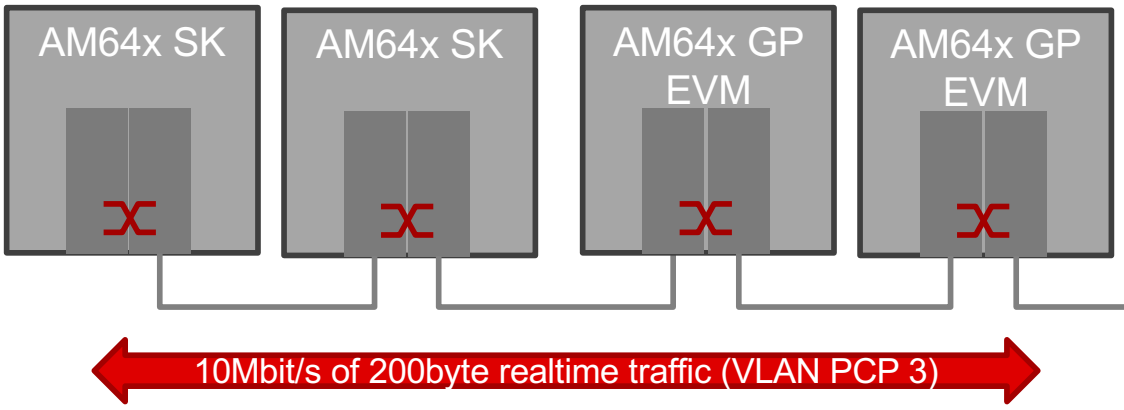

Development LAN, Linux Host

TSN Support on AM64x

- Time Sensitive Networking is an umbrella term for several optional Ethernet layer 2 features for achieving realtime performance
  - See white paper https://www.ti.com/lit/spry316 for additional background
  - Support is available in several Sitara and Jacinto processors
- The key features supported on Sitara AM64x Linux (iproute2 package) are:
  - **Integrated switching including cut-thru** (Cut-thru is not a IEEE802.1 standard yet)
    - Linux SW support and line topology demo (Sitara SDK 7.3.1)
  - **Timing over packet** (802.1AS-2011, IEEE1588, gPTP) with linuxptp (ptp4l)
  - **Preemption** (IET, 802.1Qbu/802.3br) with Linux packet scheduler (ethtool, tc qdisc)
  - **Time aware shaper** (EST, 802.1Qbv) with Linux packet scheduler (tc qdisc)
- Quality of service (QoS) including TSN works by preventing interference from lower priority traffic
  - It is the realtime protocol and application that must leverage the lack of interference to meet realtime

Enhancements to Scheduled Traffic (EST), Time Aware Shaper

- A transmit port feature typically synchronized across the local area network (LAN)
- Cyclic timeslots to allow realtime traffic to meet deadlines
  - E.g every 1000µs for 125µs only motion control traffic will be sent
- Uses 8 hardware queues per transmit port
  - Can be serialized with other QoS
- Typically require an network engineering tool to configure the LAN
  - With AM64x this would be software configuring the Linux packet scheduler at each transmit port
- Requires accurate PTP to align application endpoint and switch timing
- 128 entry gate control list, each entry contains:
  - Bitmask of open gates (queues) and duration
  - The sum of the times is the cycle time

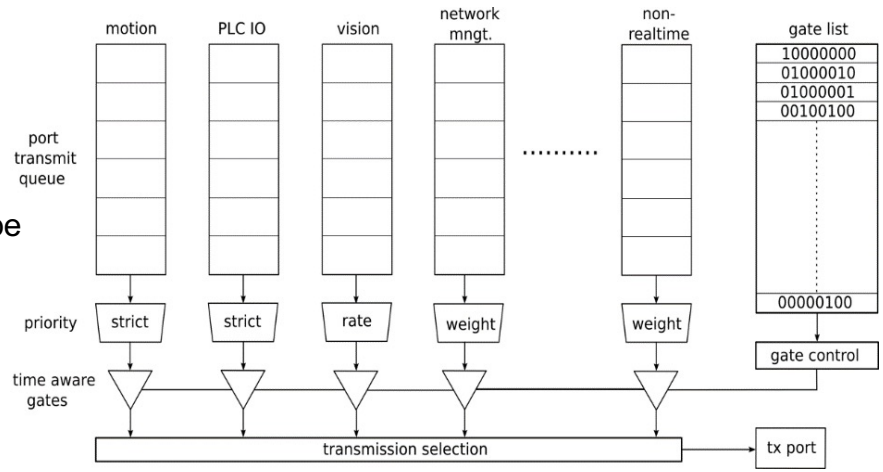

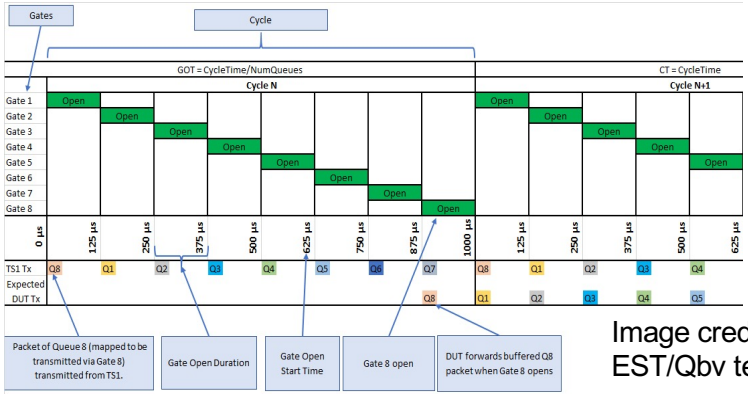

Image credit AVNU EST/Qbv test specification

Interspersing Express Traffic (IET), Preemption

- Port to port feature
- Splits MAC into two
  - Express traffic
  - Preemptable traffic
- 8 hardware queues per transmit port
  - Each queue can be express or preemptable
  - Can be serialized with other QoS
  - No nested preemption
- Does not require an network engineering tool to configure the LAN
  - LLDP messages to announce capability and minimum supported non-final fragment size
  - AM64x supports the optional further preamble level verify and respond sequence
- Does not require accurate PTP but only separates express traffic from all other traffic
- Cut-thru switching is only supported for express traffic

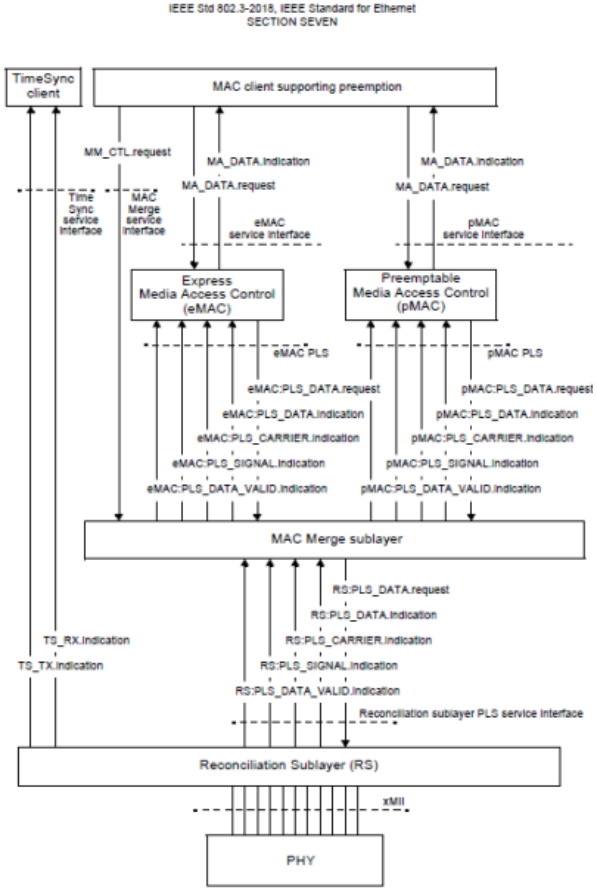

Figure 99–2—MAC Merge sublayer service interfaces diagram

Linux Packet Scheduler and TSN

- *tc* (Traffic Classes) and *qdisc* (queueing discipline)
- Maps 16 socket (SKB) priorities to queues that can optionally be offloaded to HW
  - Software shaping is relevant orders of magnitude above typical TSN system
  - Credit Based Shaper (Ethernet AVB) has been supported for years
  - CPSW support for both EST and IET is supported in 5.4 kernel
- For more general description on the topic:
  - https://tsn.readthedocs.io/qdiscs.html

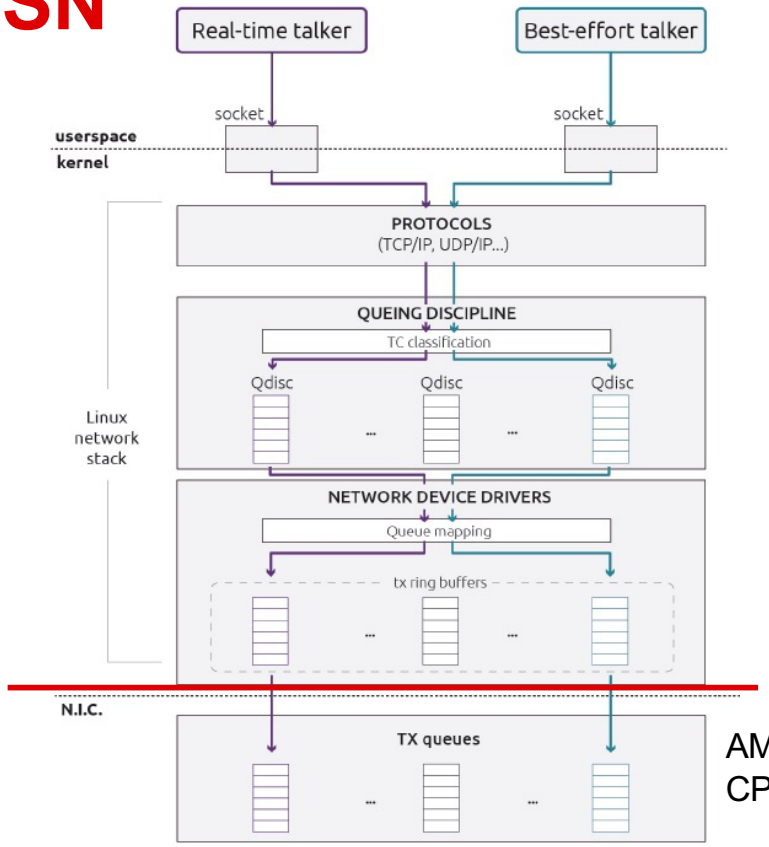

AM6x CPSW

Fig. 1: Linux networking stack transmission path

https://arxiv.org/pdf/1808.10821.pdf

Cut-Through with AM6442 (CPSW3G)

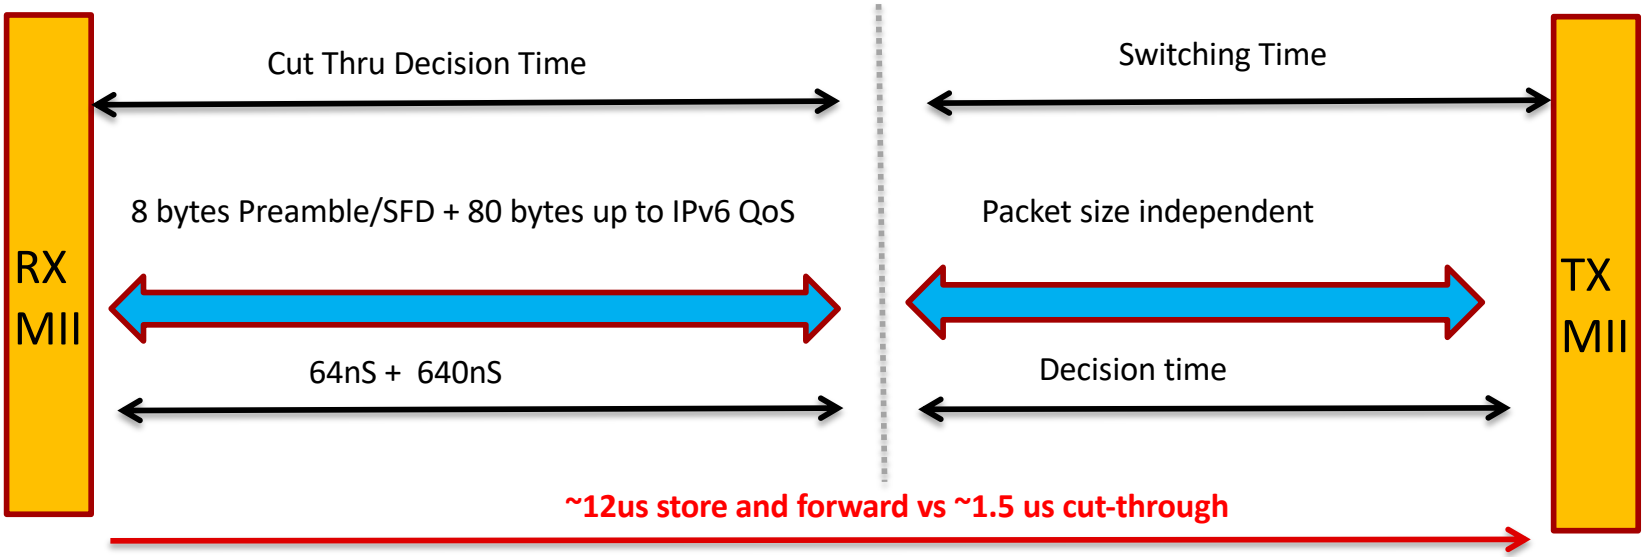

- Standard Ethernet Store and forward receives the entire frame, then forwards it
  - 1500 byte frame at 1Gbit/s takes 12us (1500 bytes times 8 ns per byte) to receive

- For Cut-Through decision is being taken after receiving 80 bytes. Lookup supports up to IPv6 QoS bits
- It takes 640 nanoseconds to receive 80 bytes at 1Gbit/s
- There is 8 byte preamble/SFD which takes 64 ns
- It takes further fixed processing to decide and start the transmission of packet.

Questions

Setup the LAN

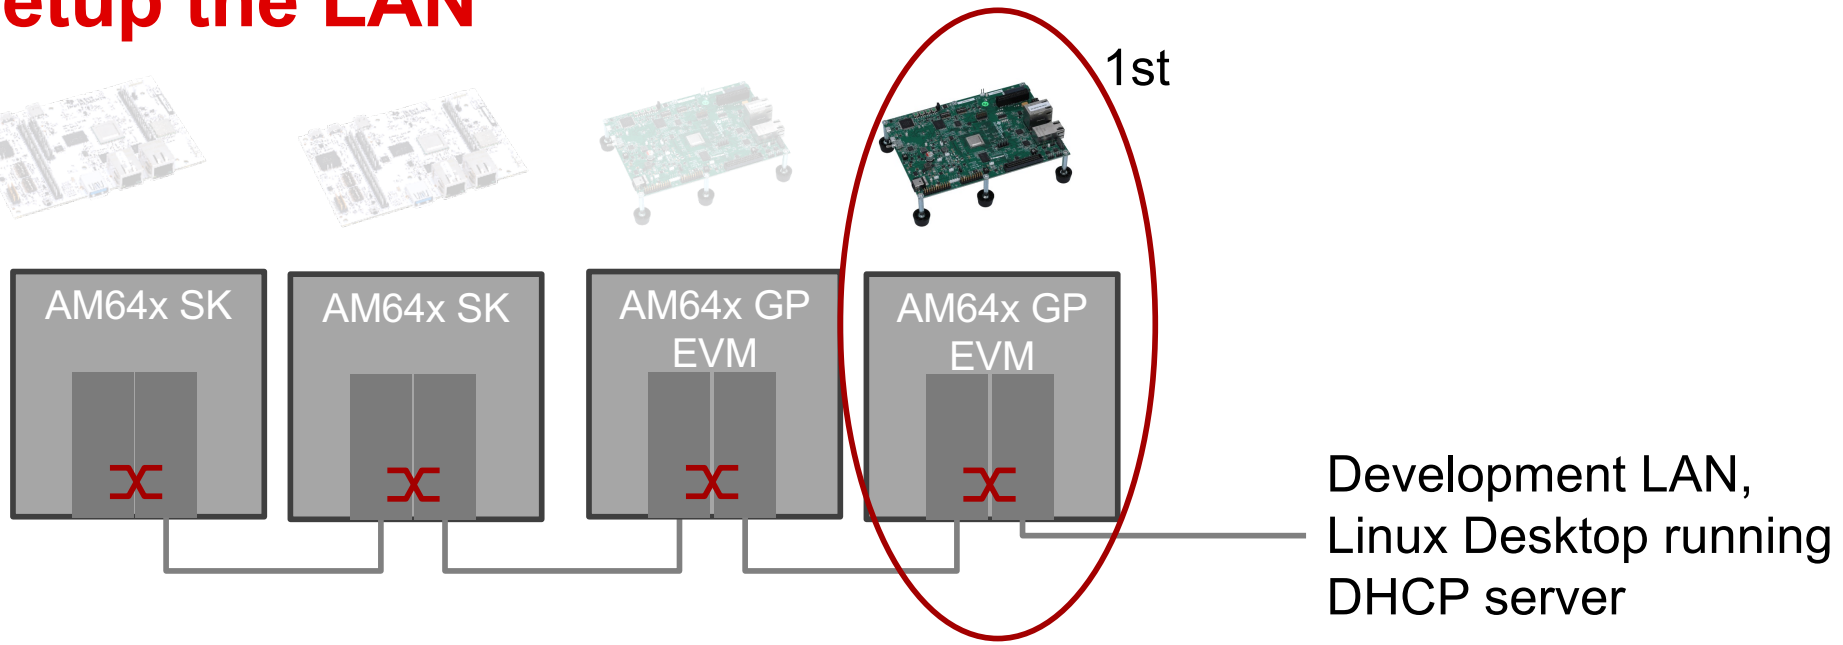

1. Connect USB console to each board (see AM64x SK Quick Start Guide)
2. Starting with the board connected to your Linux machine (on the right above)
    1. By default the AM64x EVM and SK come up in dual MAC mode
    2. Configure as cut through switch (later slide)

Setup the LAN

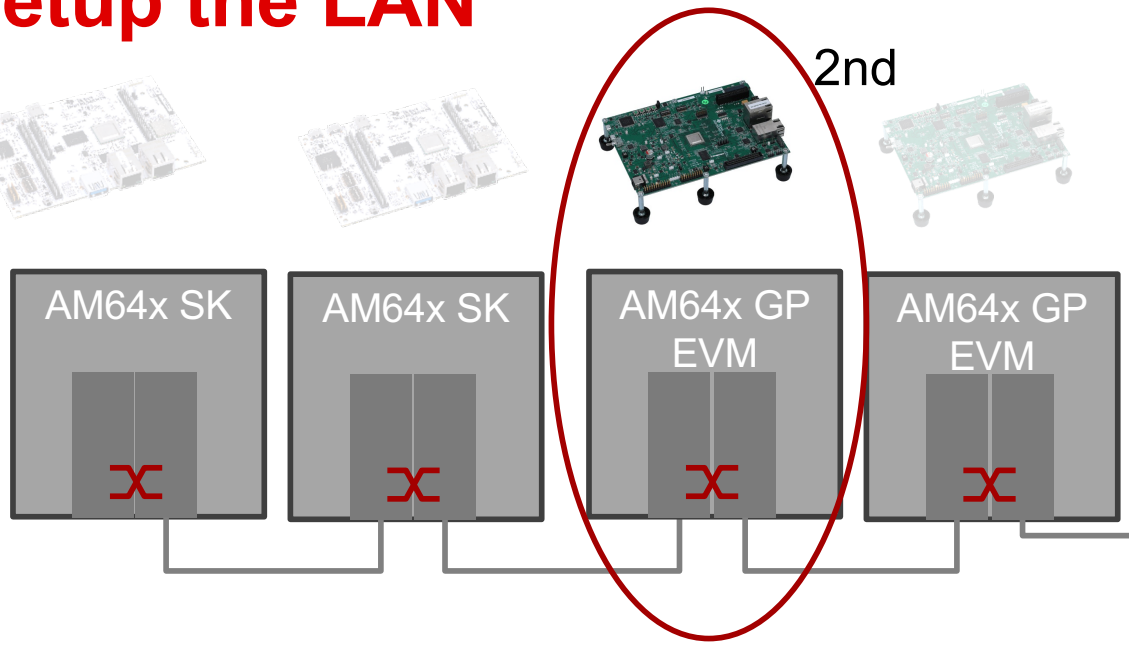

1. Connect USB console to each board (see AM64x SK Quick Start Guide)
2. Starting with the board connected to your Linux machine (on the right above)
    1. By default the AM64x EVM and SK come up in dual MAC mode
    2. Configure as cut through switch (later slide)

Setup the LAN

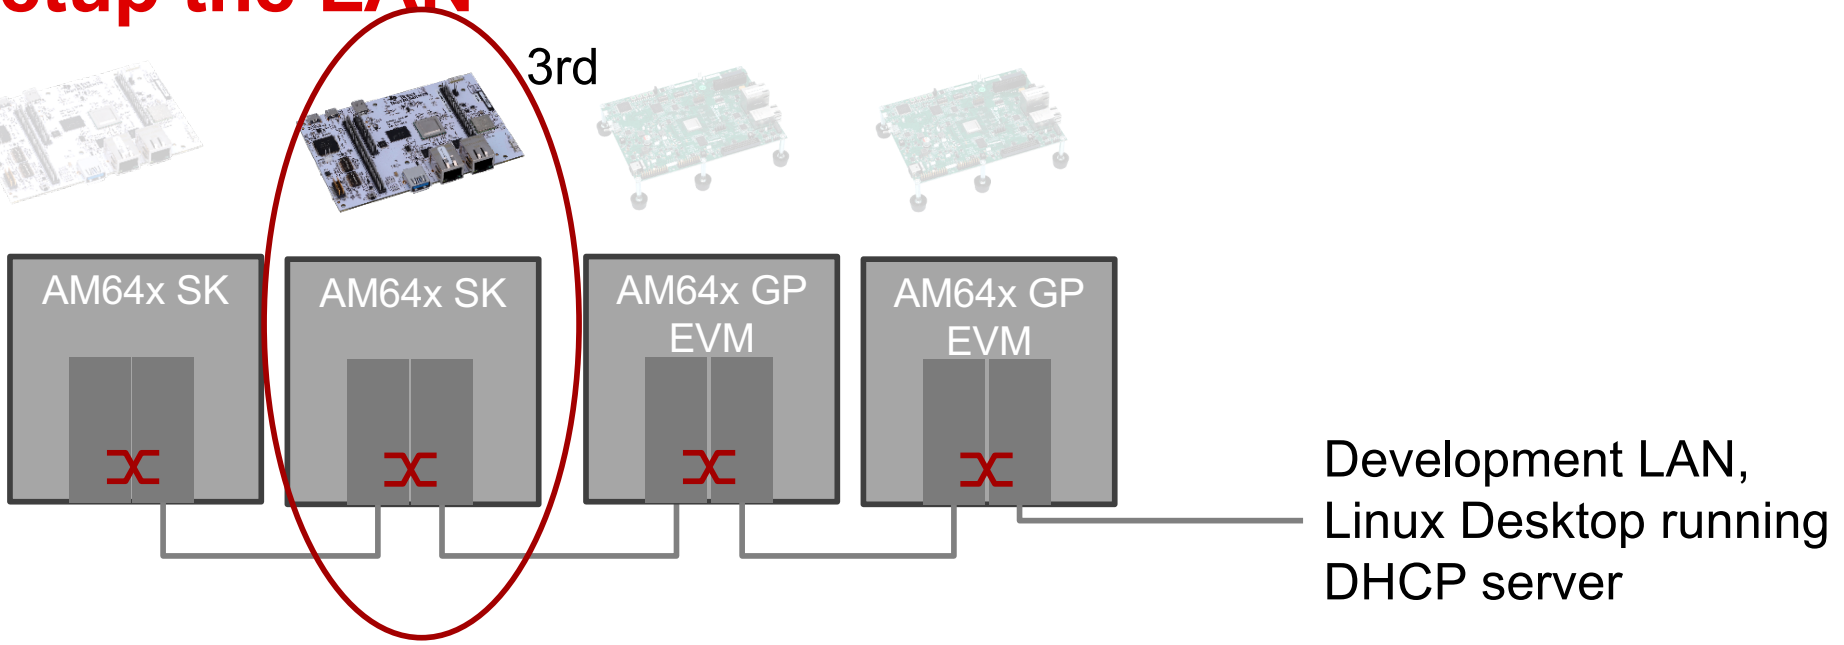

1. Connect USB console to each board (see AM64x SK Quick Start Guide)
2. Starting with the board connected to your Linux machine (on the right above)
    1. By default the AM64x EVM and SK come up in dual MAC mode
    2. Configure as cut through switch (later slide)

Setup the LAN

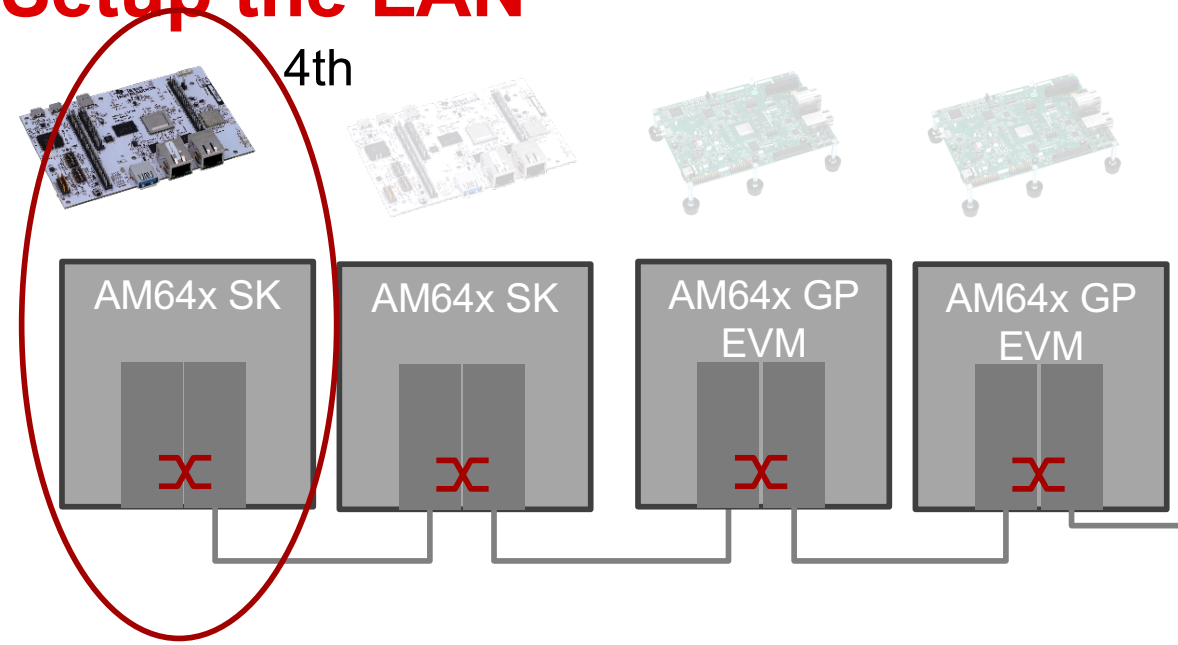

1. Connect USB console to each board (see AM64x SK Quick Start Guide)
2. Starting with the board connected to your Linux machine (on the right above)
    1. By default the AM64x EVM and SK come up in dual MAC mode
    2. Configure as cut through switch (later slide)

Script to Configure each AM64x for cut-through Switch with Preemption

```
echo "Script to initialize AM64x Ethernet ports (CPSW3G) as hardware based 802.1Q cut-through";
echo "bridge using ip and devlink commands.";
echo
echo "setting cut-thru and preemption for VLAN PCP 3 , bitmask 8 0b00001000";
echo
#take the interfaces down, by default dual MAC
ip link set dev eth0 down
ip link set dev eth1 down
# two queues one express, one non-express, and SMD level verify on
ethtool -L eth0 tx 2
ethtool -L eth1 tx 2
ethtool --set-priv-flags eth0 p0-rx-ptype-rrobin off
ethtool --set-priv-flags eth0 iet-frame-preemption on
ethtool --set-priv-flags eth0 iet-mac-verify on
ethtool --set-priv-flags eth1 p0-rx-ptype-rrobin off
ethtool --set-priv-flags eth1 iet-frame-preemption on
ethtool --set-priv-flags eth1 iet-mac-verify on
#tell the kernel hw switch to turn on switching
devlink dev param set platform/8000000.ethernet name switch_mode value true cmode runtime
#setup Q7 of 8 queues to have cut-thru on
echo 8 > /sys/kernel/debug/8000000.ethernet/Port1/cut_thru_rx_pri_mask
echo 8 > /sys/kernel/debug/8000000.ethernet/Port1/cut_thru_tx_pri_mask
echo 8 > /sys/kernel/debug/8000000.ethernet/Port2/cut_thru_rx_pri_mask
echo 8 > /sys/kernel/debug/8000000.ethernet/Port2/cut_thru_tx_pri_mask
ethtool   --set-priv-flags eth0 cut-thru  on
ethtool   --set-priv-flags eth1 cut-thru  on
```

Script to Configure each AM64x as a Switch

```
#setup a host port called br0
ip link add name br0 type bridge
ip link set dev br0 type bridge ageing_time 1000
#bring up th external Ethernet MACs, sleep is there to allow the link negotiation to go through
ip link set dev eth0 up
sleep 1
ip link set dev eth1 up
sleep 1
#connect the external MACs to the host port
ip link set dev eth0 master br0
ip link set dev eth1 master br0
#tell that the host port is part of the bridge
ip link set dev br0 type bridge stp_state 1
#bring the host port up
ip link set dev br0 up
#CPSW switch comes up in vlan aware mode, with ports blocked by default
#vlan id 1 is used for untagged traffic with VLANs, this unblock the untagged traffic to br0 por
t
bridge vlan add dev br0 vid 1 pvid untagged self
#echo spanning tree for the switch
brctl showstp br0
#get an IP address for the host port using DHCP
udhcpc -i br0
```

IEEE 802.1AS-2011 using ptp4l

- ptp4l with ptp relay at the AM64x's doing the switching
  - converges to single digit accuracy with outliers at ~10ns
  - to get the Linux clock to use this ptp time run from the command line, or set it up as a service.
- On the grandmaster (options are autoconfig, system clock, and allow system clock to be source, --transportSpecific=1 means 802.1AS domain in 1588):
  - `date -s "27 MAY 2021 14:27:10"`
  - `phc2sys -a -rr --transportSpecific=1 &`
- At the others two (options are autoconfig, system clock):
  - `phc2sys -a -r --transportSpecific=1 &`

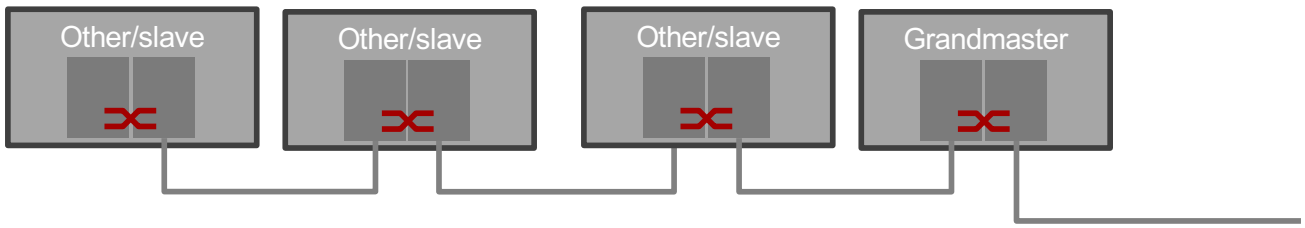

Start ptp4l

- Grandmaster AM64x first then others:
  - **`ptp4l -P -2 -H -i eth0 -i eth1 -f /etc/gPTP.cfg --step_threshold=1 -m -q -p /dev/ptp0`**
    - P means peer to peer (as opposed to end to end)
    - 2 means Ethernet level IEEE1588 (not IPv4 or IPv6)
    - H means use HW timestamps
    - i is interface
    - m is print to stdout
    - q is don't print to system log
    - p is what clock on the AM6x to use
    - gPTP.cfg is a file from https://github.com/richardcochran/linuxptp/blob/master/configs/gPTP.cfg
      - Two lines added based on the PHY datasheet on the EVM *ingressLatency 88, egressLatency 288*
    - step_threshold=1 is for faster convergence on jumps
  - On the other boards you can also check with command *date* to see Linux time has adjusted
- See https://tsn.readthedocs.io/timesync.html for more details

IEEE802.1AS-2011 running

Grandmaster

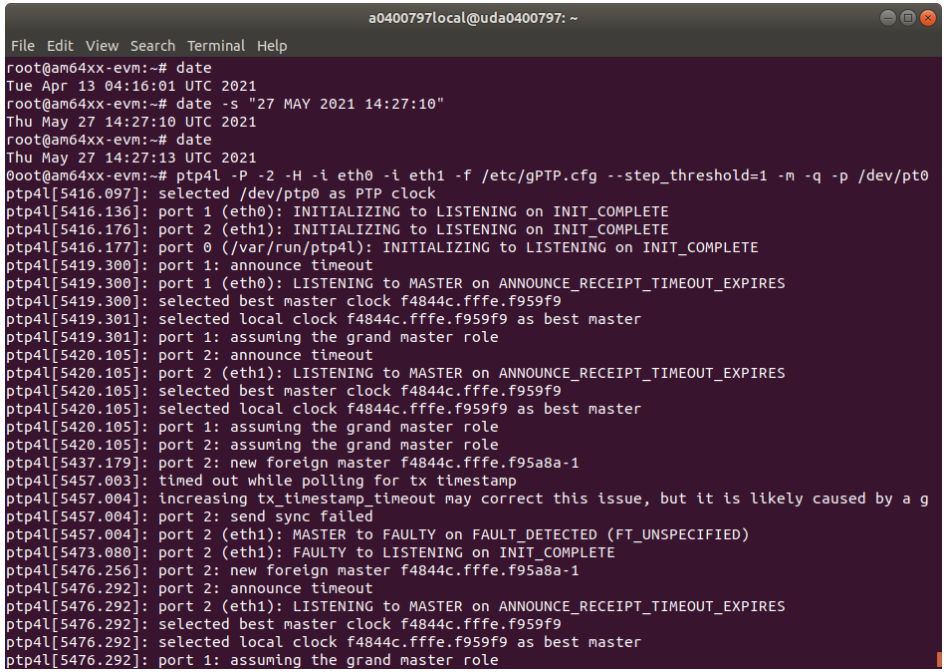

Others

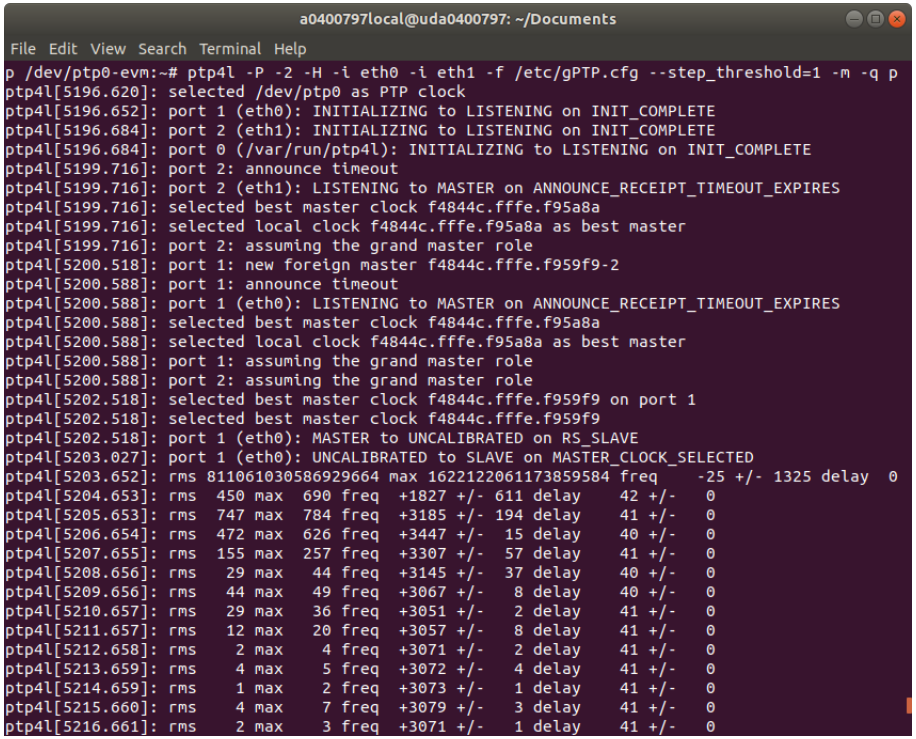

Questions

Realtime Traffic without QoS

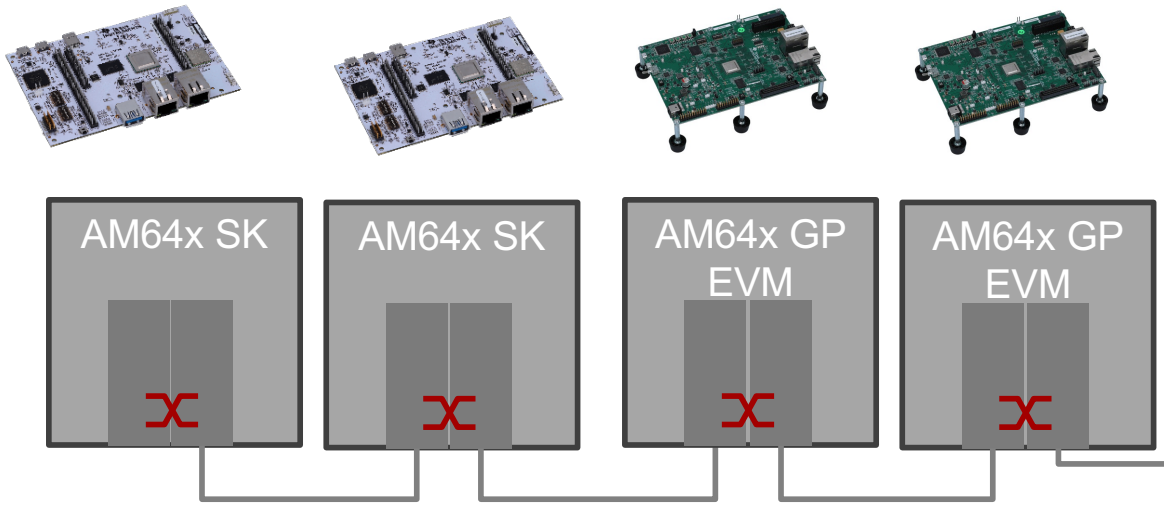

Development LAN, Linux Desktop

192.168.1.99

192.168.1.114

192.168.1.111

192.168.1.123

192.168.1.121

10Mbit/s of 200byte realtime traffic

iperf3 -u -t 0 -b 10M -c 192.168.1.121 -l 200 -p 5201

iperf3 -s -p 5201

TCP throughput test

netperf -H 192.168.1.123

Realtime Traffic without QoS

Realtime (iperf3)

Background throughput test (netperf)

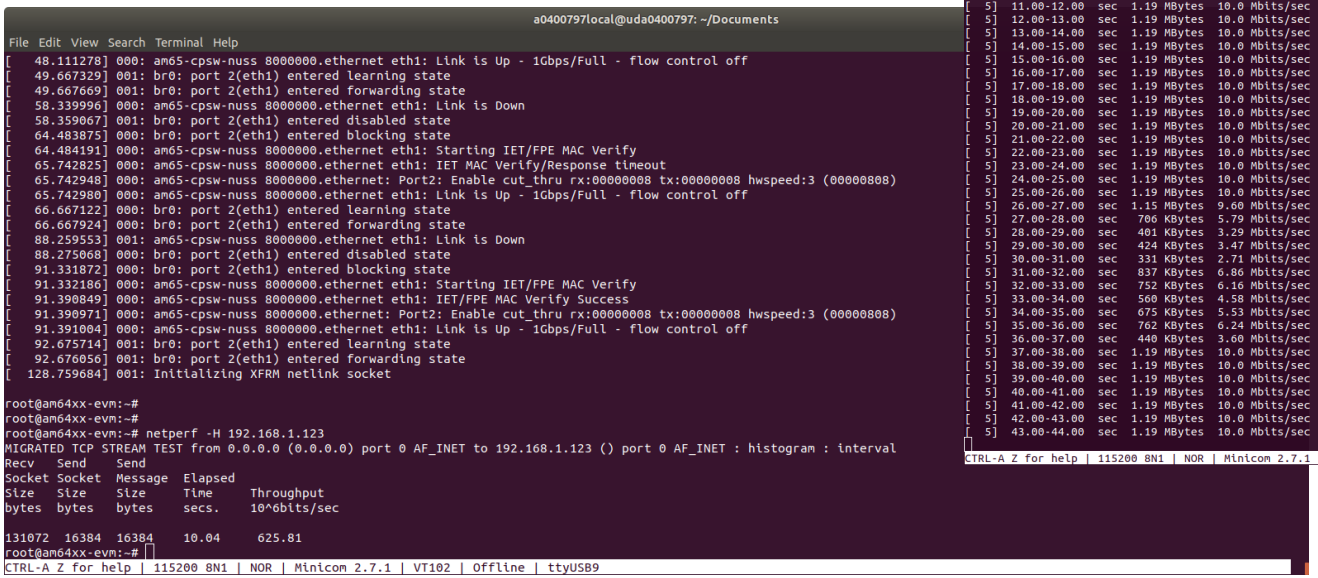

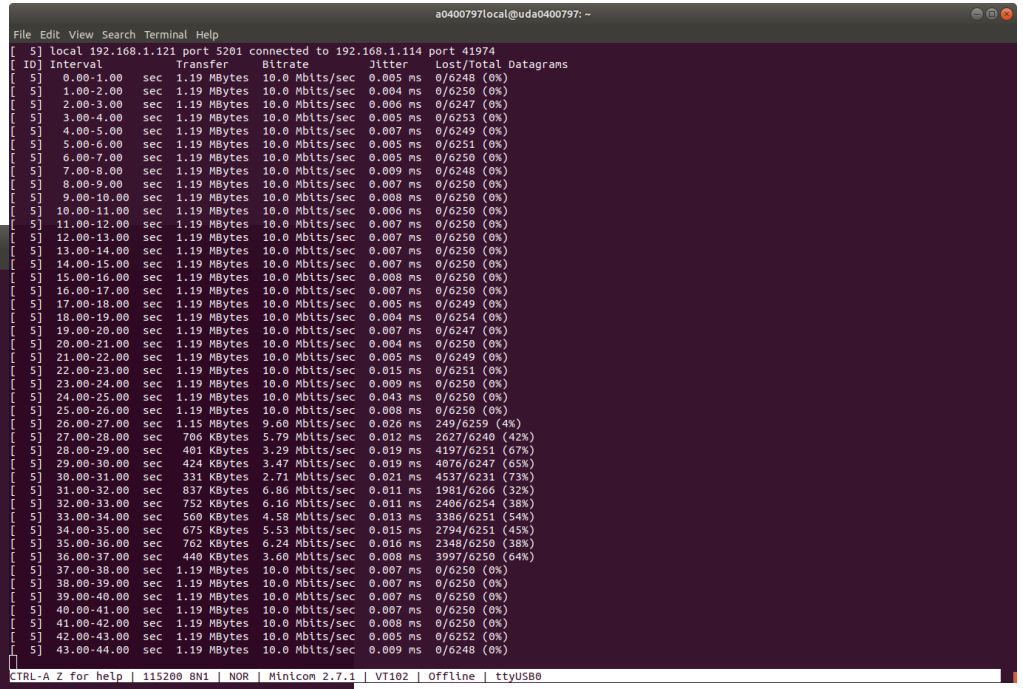

Realtime Traffic with Preemption and Cut-through

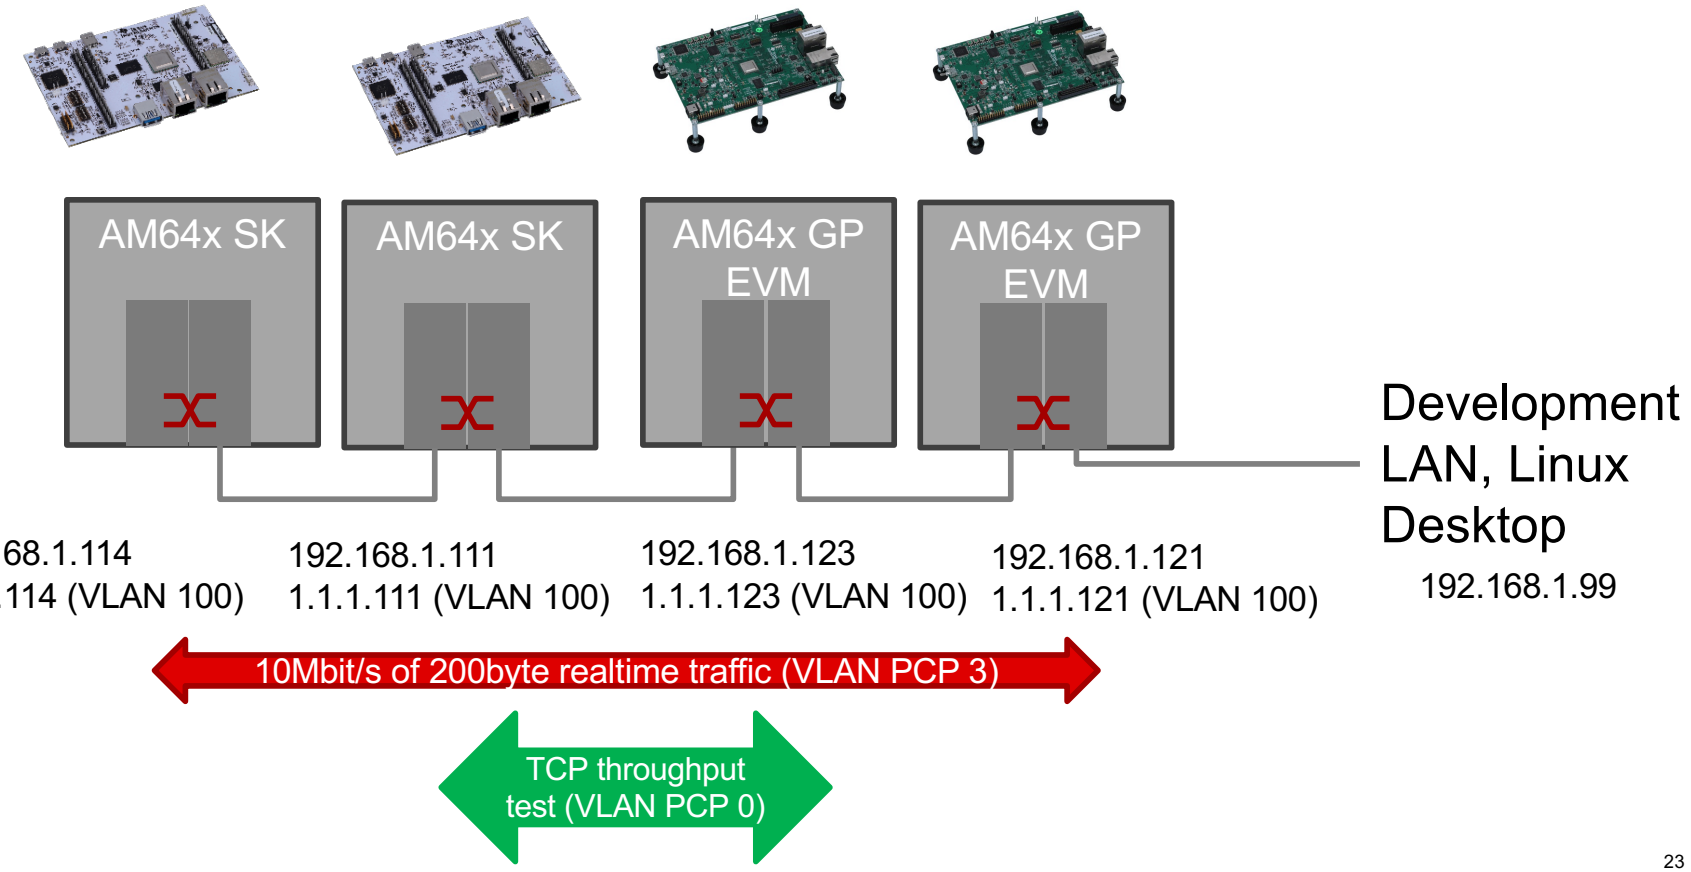

Use VLAN priority as TSN Stream Identifier

- Using VLAN priority (PCP) is perhaps the simplest usage of TSN
- Tools like iperf3 and netperf will use default priority (SO_PRIORITY) of 0
  - Map the devices with realtime so default socket uses VLAN PCP=3 which we made express and cut-through
  - In the middle two leave default as VLAN PCP=0
- Create a VLAN with ID 100 across the 4 devices

```
# middle two all traffic uses default
ip link add link br0 name br0.100 type vlan id 100
sleep 1
bridge vlan add dev br0 vid 100 self
bridge vlan add dev eth0 vid 100
bridge vlan add dev eth1 vid 100
ip link set dev br0.100 up
```

```
# endpoint two
ip link add link br0 name br0.100 type vlan id 100 egress-qos-map 0:3
sleep 1
bridge vlan add dev br0 vid 100 self
bridge vlan add dev eth0 vid 100
bridge vlan add dev eth1 vid 100
ip link set dev br0.100 up
```

- Give each VLAN device static IP address 1.1.1.x (e.g. where x is the same as the DHCP address)
  - ip addr add 1.1.1.114/24 dev br0.100
- In a real realtime application (not iperf3) you would use SO_PRIORITY other than default and appropriate mapping to TSN stream such as VLAN priority

Realtime Traffic with Cut-through and Preemeption

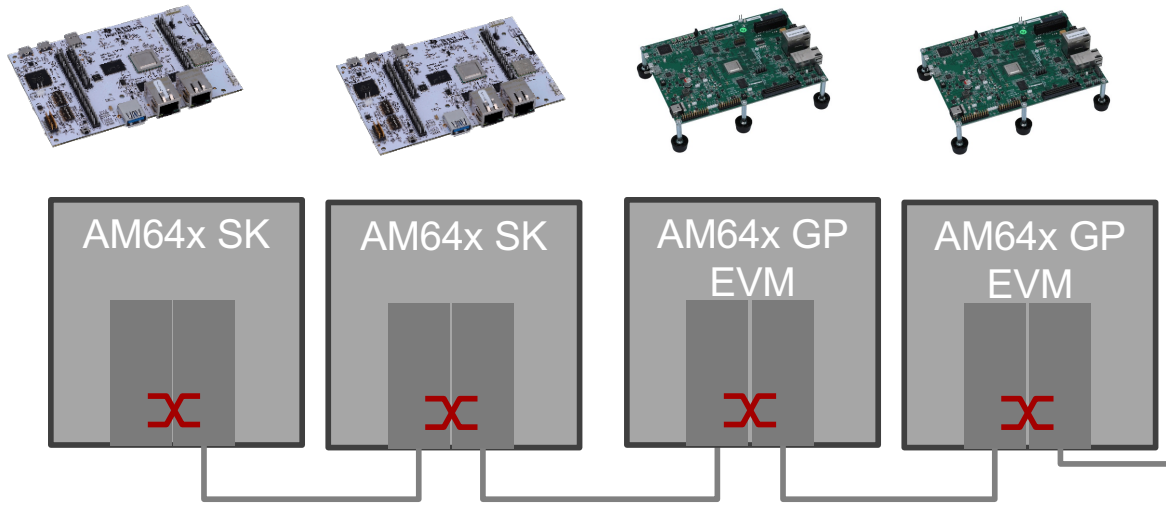

Development LAN, Linux Desktop
192.168.1.99

1.1.1.114 (VLAN 100) 1.1.1.111 (VLAN 100) 1.1.1.123 (VLAN 100) 1.1.1.121 (VLAN 100)

10Mbit/s of 200byte realtime traffic (VLAN PCP 3)

iperf3 -u -t 0 -b 10M -c 1.1.1.121 -l 200 -p 5201

iperf3 -s -p 5201

TCP throughput test V(LAN PCP 0)

netperf -H 1.1.1.123

Realtime Traffic with QoS

Realtime (iperf3) using VLAN PCP=3 is unaffected by other traffic

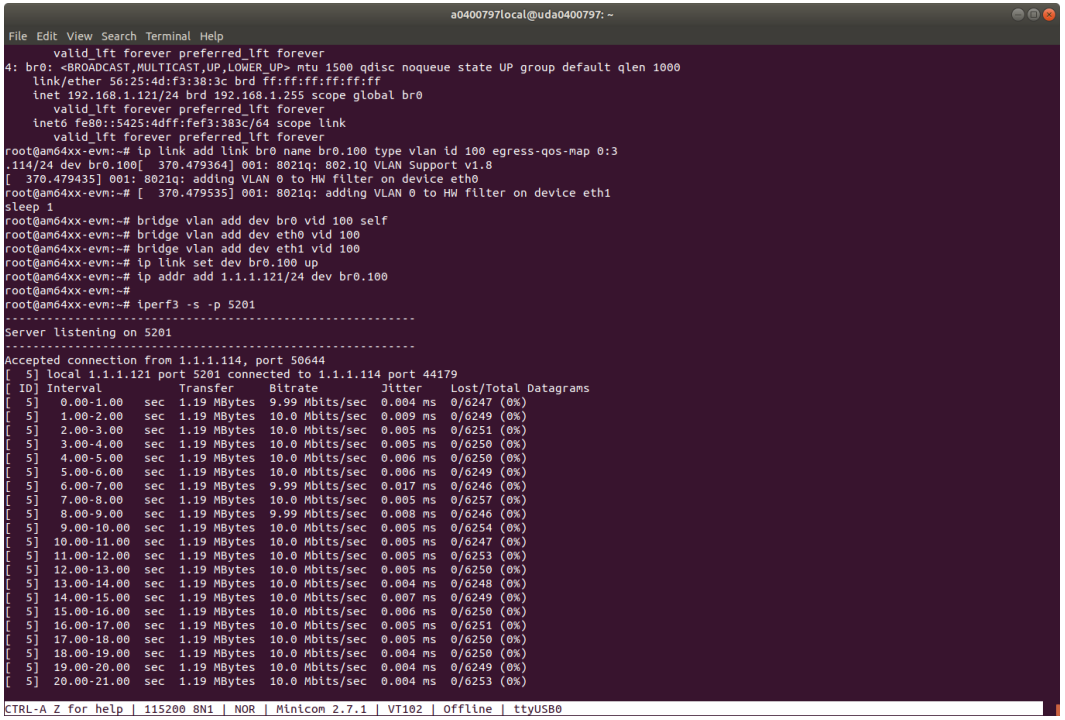

- See statistics for preemption for the background traffic using:
  - ethtool -S eth0 | grep iet
- And cut-through using:
  - ethtool -S eth1 | grep coll
  - The statistics couners for collisions have been repurposed for cut-through
    - tx_collision_frames: Tx Cut Thru with and without delay (full-duplex)
    - tx_single_coll_frames: Tx Store and Forward (full-duplex)
    - tx_mult_coll_frames: Rx Cut Thru with no delay (full-duplex)
    - tx_excessive_collisions: Rx Cut Thru with delay (full-duplex)
    - tx_late_collisions: Rx Store and Forward (full-duplex)

What about latency and EST (802.1Qbv)

- For EST (802.1Qbv) configuring the schedule across multiple switches requires an engineering tool and configuration stack
  - Building blocks for configuring an EST schedule documented at Sitara SDK EST documentation and Linux TSN documentation on EST
  - Includes endpoint only demo with tcpdump based packet capture and wireshark analysis
- For preemption and cut-through
  - A real realtime application (not iperf3) that for example toggles an IO pin and use an external analyzer
  - Microsecond level accuracy at 1Gbit/s, slowing down the interface to 100M or 10M using ethtool moves the benefit of the features to hundreds microseconds

Summary

- AM64x TSN features
- Demo
  - Configure cut-through switching with preemption
  - IEEE802.1AS-2011 using ptp4l
  - Express and background traffic
- Latency

- Questions and Answers

Resources

- Slides and videos will be posted here:
  - E2E Post
- AM64x GP EVM
- AM64x SK EVM
- AM64x Product Page
- Previous Sitara Webinar on Comms and Real-time
  - Includes Presentation and Video
  - Dives further into the Software Architecture of the Web Server Application
- Previous Sitara Webinar on general Linux Networking
  - Includes Presentation and Video coming soon
  - Covers wireless and wired switching/bridging and support topologies