



OPENFABRICS  
ALLIANCE

13<sup>th</sup> ANNUAL WORKSHOP 2017

# Asynchronous Peer-to-Peer Device Communication

Feras Daoud, Leon Romanovsky

[ 28 March, 2017 ]



# Agenda

Peer-to-Peer communication

PeerDirect technology

PeerDirect and PeerDirect Async

Performance

Upstream work




# Peer-to-Peer Communication

# Peer-to-Peer Communication

*“Direct data transfer between PCI-E devices without the need to use main memory as a temporary storage or use of the CPU for moving data.”*

## ■ Main advantages:


- Allow direct data transfer between devices
- Control the peers directly from other peer devices
- Accelerate transfers between different PCI-E devices
- Improve latency, system throughput, CPU utilization, energy usage
- Cut out the middleman






# PeerDirect Technology

# Timeline




# Prior To GPU Direct

- GPUs use driver-allocated pinned memory buffers for transfers
- RDMA driver use pinned buffers for zero-copy kernel-bypass communication
- It was impossible for RDMA drivers to pin memory allocated by the GPU
- Userspace needed to copy data between the GPU driver's system memory region and the RDMA memory region




# GPUDirect/GPUDirect P2P

- GPU and RDMA device share the same “pinned” buffers
- GPU copies the data to system memory
- RDMA device sends it from there
- Advantages
  - Eliminate the need to make a redundant copy in CUDA host memory
  - Eliminate CPU bandwidth and latency bottlenecks



# GPUDirect RDMA/PeerDirect


- CPU synchronizes between GPU tasks and data transfer
- HCA directly accesses GPU memory
- Advantages
  - Direct path for data exchange
  - Eliminate the need to make a redundant copy in host memory



# GPUDirect RDMA/PeerDirect CPU Utilization


```
while(fin) {  
    gpu_kernel <<<... , stream>>>(buf);  
    cudaStreamSynchronize(stream);  
    ibv_post_send(buf);  
    ibv_poll_cq(cqe);  
}
```

100% CPU Utilization



# GPUDirect Async/PeerDirect Async


- Control the HCA from the GPU
  - Performance
    - Enable batching of multiple GPU and communication tasks
    - Reduce latency
  - Reduce CPU utilization
    - Light weight CPU
    - Less power
- CPU prepares and queues compute and communication tasks on GPU
- GPU triggers communication on HCA
- HCA directly accesses GPU memory



# GPUDirect Async/PeerDirect Async

```
while(fin) {  
    gpu_kernel <<<... , stream>>>(buf);  
    gds_stream_queue_send(stream, qp, buf);  
    gds_stream_wait_cq(stream, cqe);  
}
```

No CPU in critical path



# Peer-to-Peer Evolution

## GPUDirect

- Eliminate the need to make a redundant copy in CUDA host memory
- Eliminate CPU bandwidth and latency bottlenecks

## PeerDirect

- Eliminate the need to make a redundant copy in host memory
- Direct path for data exchange

## PeerDirect Async

- Control RDMA device from the GPU
- Reduce CPU utilization



# PeerDirect


# PeerDirect

## How Does It Work?

- Allow `ibv_reg_mr()` to register peer memory
- Peer devices implement new kernel module – `io_peer_mem`
- Register with RDMA subsystem - `ib_register_peer_memory_client()`
- `io_peer_mem` implements the following callbacks :
  - `acquire()` – detects whether a virtual memory range belongs to the peer
  - `get_pages()` – asks the peer for the physical memory addresses matching the memory region
  - `dma_map()` – requests the bus addresses for the memory region
  - Matching callbacks for release: `dma_unmap()`, `put_pages()` and `release()`

# PeerDirect

## Memory Region Registration





# PeerDirect Async

# PeerDirect Async

## How Does It Work?

- Allows for peer devices to control the network card
  - RDMA NIC provides a bytecode sequence to the peer
  - Peer device executes bytecode to trigger sends or detect completions
- PeerDirect Async uses dedicated QPs and CQs
- PeerDirect Async operations
  - `Ibv_post_send()` on a PeerDirect Async QP queues a set of operations to be triggered by peer
  - `ibv_peer_commit_qp()` – Obtain bytecode for committing pending WQEs for execution
  - `ibv_peer_peek_cq()` – Obtain bytecode for detecting a certain number of completions
- Device agnostic
  - An network card that exports bytecode operations for `post_send` and `poll_cq`
  - Any peer device that can execute the byte code
    - GPUs, FPGAs, Storage controllers, etc.


# Transmit Operation

*Create a QP ->*

*Mark it for PeerDirect Async ->*

*Associate it with the peer*

1. Post work requests using `ibv_post_send()`
  - Doorbell is not ringed
2. Use `ibv_peer_commit_qp()` to get bytecode for committing all WQEs currently posted to the send work queue
3. Queue the translated bytecode operations on the peer
4. Peer executes the operations after generating outgoing data




# Completion Handling

*Create a CQ ->*

*Mark it for PeerDirect Async ->*

*Associate it with the peer*


1. Use `ibv_peer_peek_cq()` to get bytecode for peeking a CQ for a specific number of completions
2. Queue the translated operations on the peer before the operations that use the received data
3. Synchronize the CPU with the peer to insure that all the operations has ended
4. Use `ibv_poll_cq()` to consume the completion entries





# Performance


# Performance mode



[\*] modified ud\_pingpong test: recv+GPU kernel+send on each side.

2 nodes: Ivy Bridge Xeon + K40 + Connect-IB + MLNX switch, 10000 iterations, message size: 128B, batch size: 20

# Economy Mode



[\*] modified ud\_pingpong test, HW same as in previous slide



# Upstream Work

# Peer-to-Peer – Upstream Proposals

## ■ Peer-to-Peer DMA

- Mapping DMA addresses of PCI device to IOVA of other device

## ■ ZONE\_DEVICE

- Extend ZONE\_DEVICE functionality to memory not cached by CPU

## ■ RDMA extension to DMA-BUF

- Allow memory region create from DMA-BUF file handle

## ■ IOPMEM

- A block device for PCI-E memory

## ■ Heterogeneous Memory Management (HMM)

- Common address space will allow migration of memory between devices



OPENFABRICS  
ALLIANCE

13<sup>th</sup> ANNUAL WORKSHOP 2017

**THANK YOU**

Feras Daoud, Leon Romanovsky

